

In [145]:

```
import datashader as ds
import datashader.transfer_functions as tf
import datashader.glyphs
from datashader import reductions
from datashader.core import bypixel
from datashader.utils import lnglat_to_meters as webm, export_image
from datashader.colors import colormap_select, Greys9, viridis, inferno
import copy

from pyproj import Proj, transform
import numpy as np
import pandas as pd
import urllib
import json
import datetime
import colorlover as cl

import plotly.offline as py
import plotly.graph_objs as go
from plotly import tools

from shapely.geometry import Point, Polygon, shape
# In order to get shapely, you'll need to run [pip install shapely.geometry] from your terminal

from functools import partial

from IPython.display import GeoJSON

py.init_notebook_mode()
```

For module 2 we'll be looking at techniques for dealing with big data. In particular binning strategies and the datashader library (which possibly proves we'll never need to bin large data for visualization ever again.)

To demonstrate these concepts we'll be looking at the PLUTO dataset put out by New York City's department of city planning. PLUTO contains data about every tax lot in New York City.

PLUTO data can be downloaded from [here \(https://www1.nyc.gov/site/planning/data-maps/open-data/dwn-pluto-mappluto.page\)](https://www1.nyc.gov/site/planning/data-maps/open-data/dwn-pluto-mappluto.page). Unzip them to the same directory as this notebook, and you should be able to read them in using this (or very similar) code. Also take note of the data dictionary, it'll come in handy for this assignment.

In [146]:

```
# Code to read in v17, column names have been updated (without upper case letters) for v18

# bk = pd.read_csv('PLUTO17v1.1/BK2017V11.csv')
# bx = pd.read_csv('PLUTO17v1.1/BX2017V11.csv')
# mn = pd.read_csv('PLUTO17v1.1/MN2017V11.csv')
# qn = pd.read_csv('PLUTO17v1.1/QN2017V11.csv')
# si = pd.read_csv('PLUTO17v1.1/SI2017V11.csv')

# ny = pd.concat([bk, bx, mn, qn, si], ignore_index=True)

ny = pd.read_csv('nyc_pluto_20v1_csv/pluto_20v1.csv')

# Getting rid of some outliers
ny = ny[(ny['yearbuilt'] > 1850) & (ny['yearbuilt'] < 2020) & (ny['numfloors'] != 0)]
```

```
/Users/bobo/opt/anaconda3/lib/python3.7/site-packages/IPython/core/interactiveshell.py:3058: DtypeWarning:
Columns (17,18,20,22) have mixed types. Specify dtype option on import or set low_memory=False.
```

Columns (17,18,20,22) have mixed types. Specify dtype option on import or set low\_memory=False.

I'll also do some prep for the geographic component of this data, which we'll be relying on for datashader.

You're not required to know how I'm retrieving the latitude and longitude here, but for those interested: this dataset uses a flat x-y projection (assuming for a small enough area that the world is flat for easier calculations), and this needs to be projected back to traditional latitude and longitude.

In [147]:

```
# wgs84 = Proj("+proj=longlat +ellps=GRS80 +datum=NAD83 +no_defs
")
# nyli = Proj("+proj=lcc +lat_1=40.66666666666666 +lat_2=41.0333
3333333333 +lat_0=40.16666666666666 +lon_0=-74 +x_0=300000 +y_0=
0 +ellps=GRS80 +datum=NAD83 +to_meter=0.3048006096012192 +no_def
s")
# ny['xcoord'] = 0.3048*ny['xcoord']
# ny['ycoord'] = 0.3048*ny['ycoord']
# ny['lon'], ny['lat'] = transform(nyli, wgs84, ny['xcoord'].val
ues, ny['ycoord'].values)

# ny = ny[(ny['lon'] < -60) & (ny['lon'] > -100) & (ny['lat'] <
60) & (ny['lat'] > 20)]

#Defining some helper functions for DataShader
background = "black"
export = partial(export_image, background = background, export_p
ath="export")
cm = partial(colormap_select, reverse=(background!="black"))
```

# Part 1: Binning and Aggregation

Binning is a common strategy for visualizing large datasets. Binning is inherent to a few types of visualizations, such as histograms and [2D histograms](https://plot.ly/python/2D-Histogram/) (<https://plot.ly/python/2D-Histogram/>) (also check out their close relatives: [2D density plots](https://plot.ly/python/2d-density-plots/) (<https://plot.ly/python/2d-density-plots/>) and the more general form: [heatmaps](https://plot.ly/python/heatmaps/) (<https://plot.ly/python/heatmaps/>)).

While these visualization types explicitly include binning, any type of visualization used with aggregated data can be looked at in the same way. For example, let's say we wanted to look at building construction over time. This would be best viewed as a line graph, but we can still think of our results as being binned by year:

In [148]:

```
trace = go.Scatter(  
    # I'm choosing BBL here because I know it's a unique key.  
    x = ny.groupby('yearbuilt').count()['bbl'].index,  
    y = ny.groupby('yearbuilt').count()['bbl']  
)  
  
layout = go.Layout(  
    xaxis = dict(title = 'Year Built'),  
    yaxis = dict(title = 'Number of Lots Built')  
)  
  
fig = go.FigureWidget(data = [trace], layout = layout)  
  
fig
```

Something looks off... You're going to have to deal with this imperfect data to answer this first question.

But first: some notes on pandas. Pandas dataframes are a different beast than R dataframes, here are some tips to help you get up to speed:

---

Hello all, here are some pandas tips to help you guys through this homework:

[Indexing and Selecting \(https://pandas.pydata.org/pandas-docs/stable/indexing.html\)](https://pandas.pydata.org/pandas-docs/stable/indexing.html):  
.loc and .iloc are the analogs for base R subsetting, or filter() in dplyr

[Group By \(https://pandas.pydata.org/pandas-docs/stable/groupby.html\)](https://pandas.pydata.org/pandas-docs/stable/groupby.html): This is the pandas analog to group\_by() and the appended function the analog to summarize(). Try out a few examples of this, and display the results in Jupyter. Take note of what's happening to the indexes, you'll notice that they'll become hierarchical. I personally find this more of a burden than a help, and this sort of hierarchical indexing leads to a fundamentally different experience compared to R dataframes. Once you perform an aggregation, try running the resulting hierarchical dataframe through a [reset\\_index\(\)](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.reset_index.html) ([https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.reset\\_index.html](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.reset_index.html)).

[Reset\\_index \(https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.reset\\_index.html\)](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.reset_index.html): I personally find the hierarchical indexes more of a burden than a help, and this sort of hierarchical indexing leads to a fundamentally different experience compared to R dataframes. reset\_index() is a way of restoring a dataframe to a flatter index style. Grouping is where you'll notice it the most, but it's also useful when you filter data, and in a few other split-apply-combine workflows. With pandas indexes are more meaningful, so use this if you start getting unexpected results.

Indexes are more important in Pandas than in R. If you delve deeper into the using python for data science, you'll begin to see the benefits in many places (despite the personal gripes I highlighted above.) One place these indexes come in handy is with time series data. The pandas docs have a [huge section](http://pandas.pydata.org/pandas-docs/stable/timeseries.html) (<http://pandas.pydata.org/pandas-docs/stable/timeseries.html>) on datetime indexing. In particular, check out [resample](https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.resample.html) (<https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.resample.html>), which provides time series specific aggregation.

[Merging, joining, and concatenation \(https://pandas.pydata.org/pandas-docs/stable/merging.html\)](https://pandas.pydata.org/pandas-docs/stable/merging.html): There's some overlap between these different types of merges, so use this as your guide. Concat is a single function that replaces cbind and rbind in R, and the results are driven by the indexes. Read through these examples to

get a feel on how these are performed, but you will have to manage your indexes when you're using these functions. Merges are fairly similar to merges in R, similarly mapping to SQL joins.

Apply: This is explained in the "group by" section linked above. These are your analogs to the plyr library in R. Take note of the lambda syntax used here, these are anonymous functions in python. Rather than predefining a custom function, you can just define it inline using lambda.

Browse through the other sections for some other specifics, in particular reshaping and categorical data (pandas' answer to factors.) Pandas can take a while to get used to, but it is a pretty strong framework that makes more advanced functions easier once you get used to it. Rolling functions for example follow logically from the apply workflow (and led to the best google results ever when I first tried to find this out and googled "pandas rolling")

Google Wes Mckinney's book "Python for Data Analysis," which is a cookbook style intro to pandas. It's an O'Reilly book that should be pretty available out there.

---

## Question

After a few building collapses, the City of New York is going to begin investigating older buildings for safety. The city is particularly worried about buildings that were unusually tall when they were built, since best-practices for safety hadn't yet been determined. Create a graph that shows how many buildings of a certain number of floors were built in each year (note: you may want to use a log scale for the number of buildings). Find a strategy to bin buildings (It should be clear 20-29-story buildings, 30-39-story buildings, and 40-49-story buildings were first built in large numbers, but does it make sense to continue in this way as you get taller?)

In [136]:

```
ny.groupby('yearbuilt')
```


Out[136]:

```
<pandas.core.groupby.generic.DataFrameGroupBy object  
at 0x8c7a328d0>
```

In [149]:

```
# Start your answer here, inserting more cells as you go along
```

```
import matplotlib.pyplot as plt
plt.figure(figsize=(20,8))
plt.hist2d(x = ny.yearbuilt,
            y = ny.numfloors,
            bins = [(2020-1850)/10,(120-20)/10],
            range = [[1850,2020],[20,120]])
plt.xlabel('Decade Built')
plt.ylabel('Number of Floors Built')
plt.colorbar()
plt.show()
```

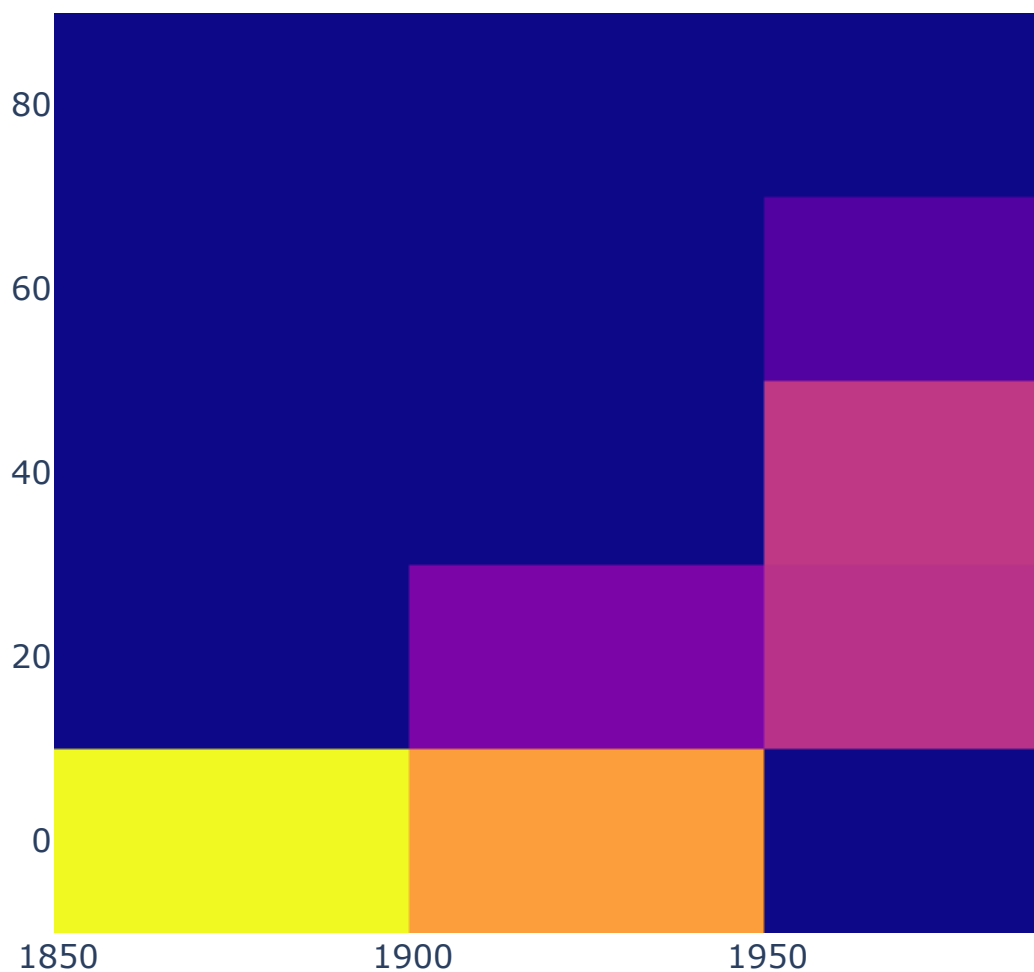


From the 2D Histogram, we see that most buildings that are taller than 20 floors were built during 1960-1970, followed by the decade between 1920 and 1930. The first unusual tall building were built in the 1890's. However, we can't really tell the different counts between 0 and 50 for buildings above 50 stories. Thus, we would create two more 2D histogram focusing on the dark blue areas in the older time of history.

In [21]:

```
# Start your answer here, inserting more cells as you go along
```

```
fig = go.Figure(go.Histogram2d(  
    x = ny.groupby('yearbuilt').count()['bbl'].index,  
    y = ny.groupby('numfloors').count()['bbl'].index  
))  
fig.show()
```






In [150]:

```
# The oldest tall buildings built before 1950
```

```
plt.figure(figsize=(8,6))
plt.hist2d(x = ny.yearbuilt,
            y = ny.numfloors,
            bins = [(1950-1890)/10,(80-20)/10],
            range = [[1890,1950],[20,80]])
plt.xlabel('Decade Built')
plt.ylabel('Number of Floors Built')
plt.colorbar()
plt.show()
```

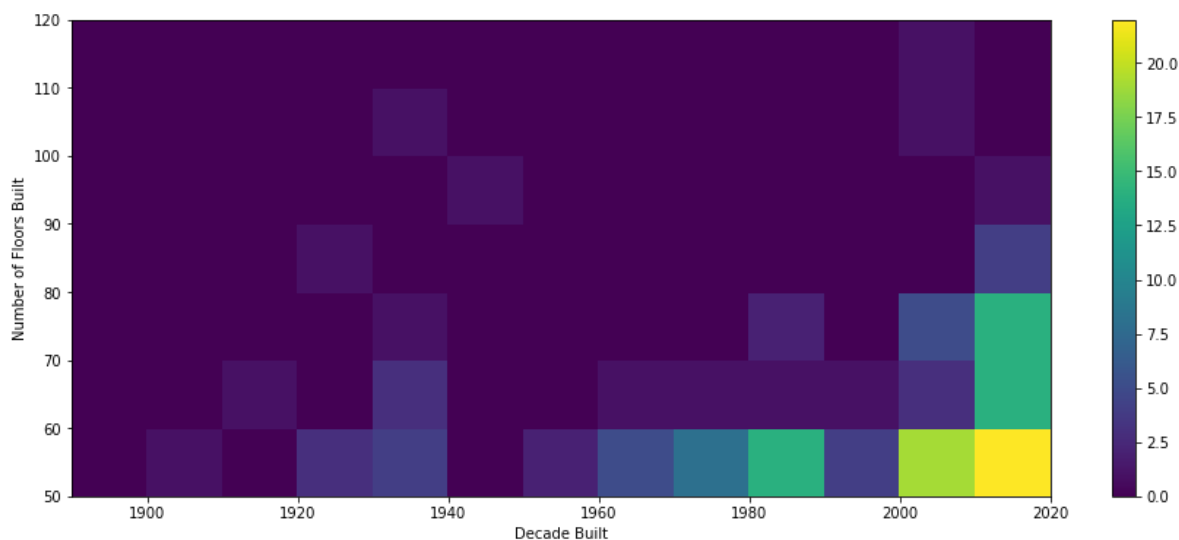


Looking at the 2D histogram above, we see that most of the unusual tall buildings were built

In [151]:

```
# Buildings that have 50 or more floors
```

```
plt.figure(figsize=(15,6))
plt.hist2d(x = ny.yearbuilt,
           y = ny.numfloors,
           bins = [(2020-1890)/10,(120-50)/10],
           range = [[1890,2020],[50,120]])
plt.xlabel('Decade Built')
plt.ylabel('Number of Floors Built')
plt.colorbar()
plt.show()
```



## Part 2: Datashader

Datashader is a library from Anaconda that does away with the need for binning data. It takes in all of your datapoints, and based on the canvas and range returns a pixel-by-pixel calculations to come up with the best representation of the data. In short, this completely eliminates the need for binning your data.

As an example, lets continue with our question above and look at a 2D histogram of YearBuilt vs NumFloors:

In [152]:

```
yearbins = 200
floorbins = 200

yearBuiltCut = pd.cut(ny['yearbuilt'], np.linspace(ny['yearbuilt'].min(), ny['yearbuilt'].max(), yearbins))
numFloorsCut = pd.cut(ny['numfloors'], np.logspace(1, np.log(ny['numfloors'].max()), floorbins))

xlabels = np.floor(np.linspace(ny['yearbuilt'].min(), ny['yearbuilt'].max(), yearbins))
ylabels = np.floor(np.logspace(1, np.log(ny['numfloors'].max()), floorbins))

fig = go.FigureWidget(
    data = [
        go.Heatmap(z = ny.groupby([numFloorsCut, yearBuiltCut])['bbl'].count().unstack().fillna(0).values,
                    colorscale = 'Greens', x = xlabels, y = ylabels)
    ]
)

fig
```


This shows us the distribution, but it's subject to some biases discussed in the Anaconda notebook [Plotting Perils](https://anaconda.org/jbednar/plotting_pitfalls/notebook) ([https://anaconda.org/jbednar/plotting\\_pitfalls/notebook](https://anaconda.org/jbednar/plotting_pitfalls/notebook)).

Here is what the same plot would look like in datashader:

In [153]:

```
cvs = ds.Canvas(800, 500, x_range = (ny['yearbuilt'].min(), ny['yearbuilt'].max()),  
                y_range = (ny['numfloors'].min(), ny['numfloors'].max()))  
agg = cvs.points(ny, 'yearbuilt', 'numfloors')  
view = tf.shade(agg, cmap = cm(Greys9), how='log')  
export(tf.spread(view, px=2), 'yearvsnumfloors')
```

Out[153]:



That's technically just a scatterplot, but the points are smartly placed and colored to mimic what one gets in a heatmap. Based on the pixel size, it will either display individual points, or will color the points of denser regions.

Datashader really shines when looking at geographic information. Here are the latitudes and longitudes of our dataset plotted out, giving us a map of the city colored by density of structures:

In [154]:

```
NewYorkCity = (( 913164.0, 1067279.0), (120966.0, 272275.0))  
cvs = ds.Canvas(700, 700, *NewYorkCity)  
agg = cvs.points(ny, 'xcoord', 'ycoord')  
view = tf.shade(agg, cmap = cm.inferno, how='log')  
export(tf.spread(view, px=2), 'firery')
```

Out[154]:



Interestingly, since we're looking at structures, the large buildings of Manhattan show up as less dense on the map. The densest areas measured by number of lots would be single or multi family townhomes.

Unfortunately, Datashader doesn't have the best documentation. Browse through the examples from their [github repo](https://github.com/bokeh/datashader/tree/master/examples) (<https://github.com/bokeh/datashader/tree/master/examples>). I would focus on the [visualization pipeline](https://anaconda.org/jbednar/pipeline/notebook) (<https://anaconda.org/jbednar/pipeline/notebook>) and the [US Census](https://anaconda.org/jbednar/census/notebook) (<https://anaconda.org/jbednar/census/notebook>) Example for the question below. Feel free to use my samples as templates as well when you work on this problem.

## Question

You work for a real estate developer and are researching underbuilt areas of the city. After looking in the [Pluto data dictionary](https://www1.nyc.gov/assets/planning/download/pdf/data-maps/open-data/pluto_datadictionary.pdf?v=17v1_1) ([https://www1.nyc.gov/assets/planning/download/pdf/data-maps/open-data/pluto\\_datadictionary.pdf?v=17v1\\_1](https://www1.nyc.gov/assets/planning/download/pdf/data-maps/open-data/pluto_datadictionary.pdf?v=17v1_1)), you've discovered that all tax assessments consist of two parts: The assessment of the land and assessment of the structure. You reason that there should be a correlation between these two values: more valuable land will have more valuable structures on them (more valuable in this case refers not just to a mansion vs a bungalow, but an apartment tower vs a single family home). Deviations from the norm could represent underbuilt or overbuilt areas of the city. You also recently read a really cool blog post about [bivariate choropleth maps](http://www.joshuastevens.net/cartography/make-a-bivariate-choropleth-map/) (<http://www.joshuastevens.net/cartography/make-a-bivariate-choropleth-map/>), and think the technique could be used for this problem.

Datashader is really cool, but it's not that great at labeling your visualization. Don't worry about providing a legend, but provide a quick explanation as to which areas of the city are overbuilt, which areas are underbuilt, and which areas are built in a way that's properly correlated with their land value.

In [161]:

```
ny.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 811684 entries, 0 to 859171
```

```
Data columns (total 99 columns):
```

borough	811684	non-null	object
block	811684	non-null	int64
lot	811684	non-null	int64
cd	811384	non-null	float64
ct2010	811384	non-null	float64
cb2010	811384	non-null	float64
schooldist	811341	non-null	float64
council	811384	non-null	float64
zipcode	811328	non-null	float64
firecomp	811339	non-null	object
policeprct	811341	non-null	float64
healtharea	811341	non-null	float64
sanitboro	811321	non-null	float64
sanitsub	811252	non-null	object
address	811684	non-null	object
zonedist1	811188	non-null	object
zonedist2	17421	non-null	object
zonedist3	119	non-null	object
zonedist4	4	non-null	object
overlay1	69789	non-null	object
overlay2	152	non-null	object
spdist1	93003	non-null	object
spdist2	64	non-null	object
spdist3	0	non-null	float64
ltdheight	2217	non-null	object
splitzone	811188	non-null	object
bldgclass	811684	non-null	object
landuse	811027	non-null	float64
easements	811684	non-null	float64
ownertype	16613	non-null	object
ownername	811663	non-null	object
lotarea	811684	non-null	float64
bldgarea	811668	non-null	float64
comarea	805300	non-null	float64
resarea	805300	non-null	float64
officearea	805300	non-null	float64
retailarea	805300	non-null	float64
garagearea	805300	non-null	float64
strgearea	805300	non-null	float64
factryarea	805300	non-null	float64
otherarea	805300	non-null	float64
areasource	811684	non-null	float64

numbldgs	811684	non-null	float64
numfloors	811684	non-null	float64
unitsres	811684	non-null	float64
unitstotal	811684	non-null	float64
lotfront	811684	non-null	float64
lotdepth	811684	non-null	float64
bldgfront	811684	non-null	float64
bldgdepth	811684	non-null	float64
ext	790509	non-null	object
proxcode	811684	non-null	float64
irrlotcode	811684	non-null	object
lottype	811684	non-null	float64
bsmtcode	811684	non-null	float64
assessland	811684	non-null	float64
assesstot	811684	non-null	float64
exempttot	811684	non-null	float64
yearbuilt	811684	non-null	float64
yearalter1	811684	non-null	float64
yearalter2	811684	non-null	float64
histdist	28556	non-null	object
landmark	1205	non-null	object
builtfar	811668	non-null	float64
residfar	811684	non-null	float64
commfar	811684	non-null	float64
facilfar	811684	non-null	float64
borocode	811684	non-null	int64
bbl	811684	non-null	int64
condono	7944	non-null	float64
tract2010	811384	non-null	float64
xcoord	811380	non-null	float64
ycoord	811380	non-null	float64
latitude	811380	non-null	float64
longitude	811380	non-null	float64
zonemap	811192	non-null	object
zmcode	14293	non-null	object
sanborn	811097	non-null	object
taxmap	811097	non-null	float64
edesignum	0	non-null	float64
appbbl	86573	non-null	float64
appdate	86573	non-null	object
plutomapid	811684	non-null	int64
version	811684	non-null	object
sanitdistrict	811321	non-null	float64
healthcenterdistrict	811341	non-null	float64
firm07_flag	26395	non-null	float64



```
pfirm15_flag      55779 non-null float64
rpaddate          0 non-null float64
dcasdate          0 non-null float64
zoningdate        0 non-null float64
landmkdate        0 non-null float64
basempdate        0 non-null float64
masdate           0 non-null float64
polidate          0 non-null float64
edesigdate        0 non-null float64
geom              811192 non-null object
dcpedited         26815 non-null object
notes             508 non-null float64
dtypes: float64(66), int64(5), object(28)
memory usage: 619.3+ MB
```

We will use columns "assessland", "assesstot", "longitude" and "latitude" from the ny dataset

In [184]:

```
# Build the dataframe
df = ny[['assessland', 'assesstot', 'longitude', 'latitude']]

# Assessment of Structure is calculated as the difference between total assessment and land assessment
assesstruct = df['assesstot'] - df['assessland']

# Add Structure Assessment column into the dataframe
df.insert(4, "AssessStruct", assesstruct)

# Show the first few rows
df.head()
```

Out[184]:

	<b>assessland</b>	<b>assesstot</b>	<b>longitude</b>	<b>latitude</b>	<b>AssessStruct</b>
0	146250.0	350550.0	-74.007347	40.637972	204300.0
1	12240.0	78900.0	-73.846003	40.786562	66660.0
2	18120.0	34380.0	-73.926923	40.653216	16260.0
3	7680.0	24600.0	-73.925958	40.623876	16920.0
4	8160.0	29760.0	-73.926030	40.623874	21600.0

In [192]:

```
pd.options.display.float_format = '{:20,.2f}'.format # remove the scientific notation, see https://stackoverflow.com/questions/17737300/suppressing-scientific-notation-in-pandas

# Take a look at the statistics of structure assessment
df.AssessStruct.describe()
```


Out[192]:

```
count          811,684.00
mean           412,635.06
std            7,106,367.80
min              0.00
25%            23,820.00
50%            36,600.00
75%            68,940.00
max          3,924,464,620.00
Name: AssessStruct, dtype: float64
```

The mean is much higher than the 75th percentile, implying that the data is heavily skewed. We confirm this with the maximum being near 4 billion. We also see minimum being zero, which wrongly suggests that the structure has no value. Compare the two boxplots below - the first one is not normally distributed; but when we take the natural log of the values, the second boxplot looks much better.


In [242]:

```
plt.boxplot(df.AssessStruct)  
plt.show()
```



In [241]:

```
plt.boxplot(np.log(df.AssessStruct))  
plt.show()
```



In [243]:

```
# Remove the zero land assessments
df = df.loc[df['assessland'] != 0]

# Take the natural log of land assessment
df['assessland'] = np.log(df['assessland'])

# Break land assessment into 3 categories
# df = df.assign(LogLand = np.log(df['assessland']))
df = df.assign(LandValue = pd.cut(df['assessland'],
                                np.linspace(start = df['assess
land'].min(), stop = df['assessland'].max(), num = 3+1),
                                labels=["L", "M", "H"])))

df.groupby('LandValue').size()
```

Out[243]:

```
LandValue
L      418317
M      385960
H         4898
dtype: int64
```

In [244]:

```
# Remove the zero structure assessments
df = df.loc[df['AssessStruct'] != 0]

# Take the natural log of structure assessment
df['AssessStruct'] = np.log(df['AssessStruct'])

df = df.assign(StructValue = pd.cut(df['AssessStruct'],
                                   np.linspace(start = df['AssessStruct'].min(), stop = df['AssessStruct'].max(), num = 3+1),
                                   labels=["L", "M", "H"])))

df.groupby('StructValue').size()
```

Out[244]:

```
StructValue
L          18544
M          777510
H           13121
dtype: int64
```

In [253]:

```
# Combine land and structure categories into a total of 9 categories
df = df.assign(LandStructV = df['LandValue'].astype(str) + df['StructureValue'].astype(str)) # see https://stackoverflow.com/questions/19377969/combine-two-columns-of-text-in-dataframe-in-pandas-python

# Remove the NaNs
df = df.loc[df['LandStructV'] != 'Lnan']
df = df.loc[df['LandStructV'] != 'nanL']

df.groupby('LandStructV').size()
```

Out[253]:

```
LandStructV
HH          4427
HM          471
LL        13630
LM       404686
MH          8694
ML          4913
MM       372353
dtype: int64
```

In [254]:

```
# Turn LandStructV into categorical data
df['LandStructV'] = pd.Categorical(df['LandStructV'])

colorKeys = {'HH': '#3b4994', 'HM': '#5698b9', 'HL': '#5ac8c8',
             'MH': '#8c62aa', 'MM': '#a5add3', 'ML': '#ace4e4',
             'LH': '#be64ac', 'LM': '#dfb0d6', 'LL': '#e8e8e8'}

NYC = ((-74.29, -73.69), (40.49, 40.92))
cvs = ds.Canvas(700, 700, *NYC)
agg = cvs.points(df, 'longitude', 'latitude', ds.count_cat('LandStructV'))
view = tf.shade(agg, color_key = colorKeys)
export(tf.spread(view, px=1), 'Assessment')
```

Out[254]:



Purple areas represent both land and structure assessments are high, we can see mostly are in Manhattan.

Underbuilt - blue areas represent high land value and low structure value, mostly distributed among Brooklyn and Queens.

Overbuilt - pink areas represent low land value and high structure value, mostly distributed among Staten Island and upper Bronx.